

# 236609 - AI and Robotics - Fall 2024

## Lesson 2a: Three-Layered Decision Making in Robotics

---

Sarah Keren

The Taub Faculty of Computer Science  
Technion - Israel Institute of Technology

# Table of contents

1. A Running Example
2. 3 Layers of Decision Making
3. Task Planning
4. Motion Planning
5. Control

## A Running Example


---

# Example




The person: "I an hungry"


# Task Planning



The person: "I an hungry"

# Motion Planning





## 3 Layers of Decision Making

---



# Components of a Robotic Agents


Sensors




Controllers



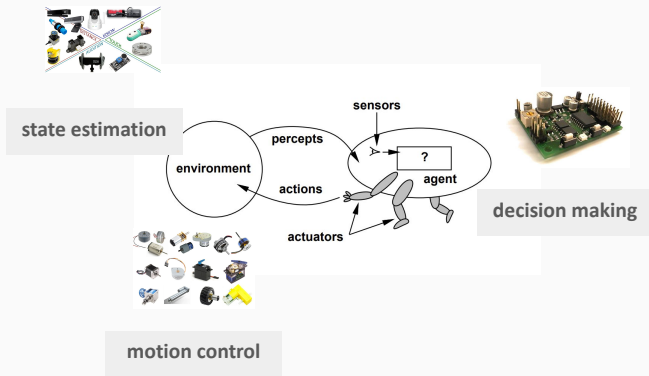
Actuators



# Autonomy



# Autonomy



# Autonomy

## State Estimation

$$\beta : \mathcal{S} \mapsto [0, 1]$$

Process incoming observations to maintain a *belief* as a probability distribution over states

## Decision Making

$$\pi : \beta \mapsto \mathcal{A}$$

$$\pi : \beta \times \mathcal{A} \mapsto [0, 1]$$


Find a policy - a mapping belief and objective into actions (probabilities)

## Motion Control

$$\dot{x}(t) = f(x(t), u(t), t) + w(t)$$

Translate actions into low-level commands (and monitor their execution)

- $x(t) \in \mathbb{R}^n$  - state vector
- $u(t) \in \mathbb{R}^m$  - control input
- $f : \mathbb{R}^n \times \mathbb{R}^m \times \mathbb{R} \rightarrow \mathbb{R}^n$  - system dynamics
- $w(t)$  - noise or disturbance



# Layered Control Architecture (LCA)

Semantic Logic Discrete Planning	<b>Decision Making</b>	Flexible and Slow
Optimization Sampling Methods Continuous Planning	<b>Trajectory Planning</b>	Intermediate
PID Control CLFs/CBFs	<b>Feedback Control</b>	Rigid and Real Time



## Top-Down Flow:

- Goals → Plans → Commands
- Like company policies becoming specific actions

## Bottom-Up Flow:

- Sensor data → Status updates → Results
- Like workers reporting to management

## Different Speeds for Different Needs


- Planning: Seconds to minutes
- Behavioral: Fraction of seconds
- Execution: Milliseconds
- Control: Microseconds
- Hardware: Nanoseconds

# Task Planning

---



# What do we need to model ?




A robot controlled by a model-based AI program

- Model of the robot
- Model of the world
- Model of possible actions
- Model of other agents

AI program reasons about the model to make decisions

# Modeling the Environment



What do we need to model ?

What do we need to model ?

## What do we need to model ?

- Sequential decision making
- Uncertainty in action outcomes
- Partial observability


## Additional considerations ?

Model needs to be compact enough to be solvable but informative enough to be useful.

# Environment: State and action space

- **State Space:**  $s \in \mathcal{S}$  is a world state (more accurately, a representation of a world state)
  - Typically, it is impossible / impractical to explicitly maintain the state space.
  - We therefore use a **factored representation** in which the state space is described via a set of variables  $\mathcal{X} = X_1, \dots, X_n$ , and a state is an assignment of a value  $X_i \in \text{Dom}(X_i)$  for each variable  $X_i$ .
- **Action space:**  $a \in \mathcal{A}$  is an action agents can perform
  - can have deterministic / non-deterministic / stochastic effects
  - may be associated with preconditions and effects.


# Examples



What are states and actions ?

# The Agent-Environment Interface

- An agent operates in the environment by taking actions.
- An agent (and its behavior) is characterized by its:
  - **Objectives** (reward and utility)
  - Ability to **observe** and **sense** the environment (observability and beliefs)



# The Agent-Environment Interface

The agent and environment interact at each of a sequence of discrete time steps  $t = 1, 2, 3, \dots$


At each time step  $t$ :

- The agent receives some representation of the environment's state -  $s_t \in \mathcal{S}$
- On that basis selects an action  $a_t \in \mathcal{A}(s)$

One time step later  $t + 1$ :

- As a consequence of its action, the agent receives a numerical reward  $r_{t+1} \in R$
- Finds itself in a new state  $s_{t+1}$

The environment and agent together thereby give rise to a sequence or **trajectory**:  $s_0, a_0, r_1, s_1, a_1, r_2, s_2, a_2, r_3, \dots$






# The Agent-Environment Interface

- **Episodes:** when the agent–environment interaction breaks naturally into sub-sequences. Each episode ends in a special state called the *terminal state*.
- **Episodic tasks:** Tasks with episodes. The time of termination  $T$  is a random variable that normally varies from episode to episode.
- **Continuing tasks:** When the agent–environment interaction does not break naturally into identifiable episodes, but goes on continually without limit.

Real-world examples of episodic vs. continuing tasks ?

# Reward



- **Reward:** a signal passed from the environment to the agent (typically referred to as **cost** when reward is negative).
- A reward  $r_t$  is a scalar feedback signal Indicates how well agent is doing at step  $t$ .
- The reward signal is your way of communicating to the agent *what* you want achieved, not *how* you want it achieved.

<https://deepmind.com/research/publications/2021/Reward-is-Enough>

## Reward (cont.)

- Often described by a reward function:
  - Depending on the setting, defined as  $\mathcal{R}(s, a, s')$ ,  $\mathcal{R}(s, a)$  or  $\mathcal{R}(s)$
  - The reward can be stationary  $\mathcal{R}(s, a, s') \in \mathbb{R}$  or non-stationary, in which case we consider the expectation
$$\mathcal{R}(s, a, s') = \mathbb{E}[r_{t+1} | S_t = s, A_t = a, S_{t+1} = s']$$
- An agent's objectives is to maximize it's **utility**  $\mathcal{U}$ 
  - the expected total reward (*return*) it receives
  - the min reward
- Utility is sometimes known as *goal* - but we will refer to a goal as a state (or state set) an agent aims to reach.

<https://deeppmind.com/research/publications/2021/Reward-is-Enough>

# Reward (by David Sliver)

- Fly stunt manoeuvres in a helicopter
  - + for following desired trajectory
  - - for crashing
- Defeat the world champion at Backgammon
  - +/- reward for winning/losing a game
- Manage an investment portfolio
  - + reward for each \$ in bank
- Control a power station
  - + reward for producing power
  - - reward for exceeding safety thresholds
- Make a humanoid robot walk
  - + reward for forward motion
  - - reward for falling over
- Play many different Atari games better than humans
  - +/- reward for increasing/decreasing score

Problems with this approach ?

# Reward

Designing a good reward function is an art.



Ball in cup image from Kober et al (2009)






Figure 2. The Competition Arena.




# Observation and Belief



- At each time step the agent receives an observation  $o_t$  that reflects the current state of the world.
- A **belief**  $\beta \in \mathcal{B}$  represents the possible world states (i.e., the states that are deemed possible by the agent).
- Induced by an **observability/ sensor function** that maps states to observations (more on this later)

# Policy (for a single agent<sup>1</sup>)

- A policy is a mapping from states to actions.
- More accurately: we seek a mapping from the agent's understanding of the state, i.e., its **belief**, to actions.
- Our objective is to find a **(sub)-optimal policy** for an agent to follow in order to achieve it's objective / maximize its **utility**.



---

<sup>1</sup>we will consider multi-agent policies in the second part of the course

# Formal definitions: Policy

- A **deterministic** policy is a mapping:
  - $\pi : \mathcal{S} \rightarrow \mathcal{A}$  from states to actions
  - $\pi : \mathcal{B} \rightarrow \mathcal{A}$  from beliefs to action
  - $\pi(s) / \pi(\beta)$  is the (single) action the agent will perform at state  $s$  or belief  $\beta$ .
- A **non-deterministic** policy is a mapping:
  - $\pi : \mathcal{S} \rightarrow \mathcal{A}^n$  from states to actions,
  - $\pi : \mathcal{B} \rightarrow \mathcal{A}^n$  from beliefs to actions
  - $\pi(s) / \pi(\beta)$  is the set of actions one of which the agent will perform at state  $s$  or belief  $\beta$ .
- A **stochastic** policy is a mapping:
  - $\pi : \mathcal{S} \times \mathcal{A} \rightarrow [0, 1]$  from states to actions
  - $\pi : \mathcal{B} \times \mathcal{A} \rightarrow [0, 1]$  from beliefs to action
  - $\pi(s, a) / \pi(\beta, a)$  is the probability the agent will perform action  $a$  at state  $s$  / belief  $\beta$ .



# Policies for Reactive Agents



## SIMPLE ROOMBA

```
IF BUMP = TRUE  
THEN Turn (random direction/amt)  
ELSE MOVE-STRAIGHT
```

<https://www.youtube.com/watch?v=7FSUtSurqA4>

# Policies for Rational Agents

- Typically cannot be described explicitly for each state and require a compact representation.
- An **optimal policy**, denoted  $\pi^*$  is *better* than or equal to all other policies.
- Since our focus is on rational agents with a clearly defined utility measure, an optimal policy maximizes the agent's utility function  $\mathcal{U}$ .



# Learning from a sample set

In supervised learning:



- A sample set:  $S = \{(x_i, y_i)_{i=1}^m\}$
- Typically used to learn for an appropriate classifier (e.g., with lowest expected error or loss) among a hypothesis class  $H$ .

Learning a policy  $\pi$ :

- A sample set:  $S = \{(s_i, a_i)_{i=1}^m\}$  or  $S = \{(\beta_i, a_i)_{i=1}^m\}$
- Typically, a parameterized policy over a factored state-space representation is used.
- Parameterized policy as a conditional probability  $\pi_\theta(a_t|\beta_t)$  or  $\pi_\theta(a_t|s_t)$  (for fully observed)

# Sample set representation of our domains ?

A sample set:  $S = \{(s_i, a_i)_{i=1}^m\}$  or  $S = \{(\beta_i, a_i)_{i=1}^m\}$



Effective for sequential decision making ?

# Propositional STRIPS

Propositional STRIPS planning task defined by a tuple

$P = \langle \mathcal{F}, I, \mathcal{A}, G, \mathcal{C} \rangle$ , where

- $\mathcal{F}$  is a set of fluents and a state  $s$  is represented by the fluents that are true in  $s$
- $I \subseteq \mathcal{F}$  is the initial state
- $G \subseteq \mathcal{F}$  represents the set of goal states, and
- $\mathcal{A}$  is a set of actions.
  - Each action is a triple  $a = \langle pre(a), add(a), del(a) \rangle$ , which represents the precondition, add, and delete lists respectively, all subsets of  $\mathcal{F}$ .
  - An action  $a$  is applicable in state  $s$  if  $pre(a) \subseteq s$ .
  - If action  $a$  is applied in state  $s$ , it results in a new state  $s' = (s \setminus del(a)) \cup add(a)$ .
- $\mathcal{C} : \mathcal{A} \rightarrow \mathbb{R}_0^+$  is a cost function that assigns each action a non-negative cost.

# Propositional STRIPS

- The objective is to find a plan  $\pi = \langle a_1, \dots, a_n \rangle$ : a sequence of actions that brings an agent from the initial state to a goal state.
- The cost  $c(\pi)$  of a plan  $\pi$  is  $\sum_{i=1}^n (C(a_i))$ .
- Often, the objective is to find an optimal solution for  $P$ : a plan  $\pi^*$  that minimizes the associated cost.
- The literature is rich with different approaches developed to solve the planning problem (Bonet and Geffner 2013): more on this later on in the course.

From [https://ai.dmi.unibas.ch/misc/tutorial\\_aaai2015/planning02.pdf](https://ai.dmi.unibas.ch/misc/tutorial_aaai2015/planning02.pdf)

# Markov Decision Process

A **Markov Decision Process**(MDP) is a tuple  $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$  where

- $\mathcal{S}$  is a finite set of states defined via a set of random variables  
 $\mathcal{X} = X_1, \dots, X_n$
- $\mathcal{A}$  is a finite set of actions
- $\mathcal{P}$  is a state transition probability matrix  
 $\mathcal{P}_{s,s'}^a = \mathcal{P}[S_{t+1} = s' | S_t = s, A_t = a]$
- $\mathcal{R}$  is a reward function,  $\mathcal{R}_s^a = \mathbb{E}[R_{t+1} | S_t = s, A_t = a]$ , and
- **optional:**  $\gamma$  is a discount factor  $\gamma \in [0, 1]$  that is used to favor immediate rewards over future rewards.

The Markov property: “The future is independent of the past given the present”.

Extensions: Infinite and continuous MDPs, partially observable MDPs, undiscounted, average reward MDPs. etc.

## Rewards (from Sutton and Barto 2018)

- Expected return for episodic tasks:

$$G_t = r_{t+1} + r_{t+2} + r_{t+3} + \dots + r_T = \sum_{k=0}^T r_{t+k+1}$$

- Expected return for continuing tasks:

$$G_t = r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots = \sum_{k=0}^{\infty} \gamma^k r_{t+k+1}$$

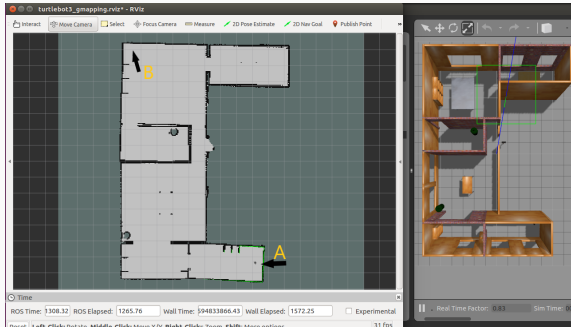
when  $\gamma$  is the discount factor.

- Returns at successive time steps are related to each other:

$$G_t = r_{t+1} + \gamma(r_{t+2} + \gamma r_{t+3} + \gamma^2 r_{t+4} + \dots) = r_{t+1} + \gamma G_{t+1}$$




# How to model our domain as MDPs ?




Is this model appropriate? Limitations? Strengths?

# Accounting for Partial Observability



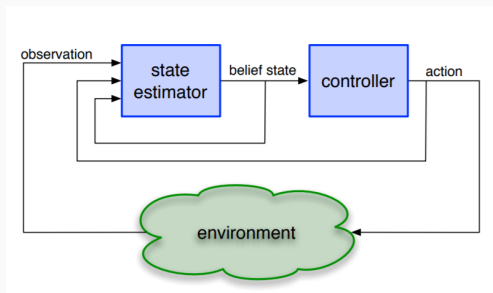
## Reminder: Observation and Belief



At each time step the agent receives an observation  $o_t$  that reflects the current state of the world.  $s$

# Beliefs and Belief Tracking

- A **belief** is a representation of the possible world states.
- In partially observable domains, we may have a **sensor model** represented as a mapping function from what is observed to the actual world state.
- The agent maintains its belief via a *state estimator* - which we will refer to as the process of **Belief Tracking**.



A belief  $\beta \in \mathcal{B}$  can have different representations:


- **Deterministic:** a direct mapping from states to beliefs  $\beta = s$  (in which case we simply talk about states)
- **Non-deterministic:**  $\beta \subseteq \mathcal{S}$  represents the set of possible world states
- **Stochastic:** the belief is a probability distribution over possible underlying world states.  $\beta : \mathcal{S} \rightarrow \mathcal{P}[\mathcal{S}]$  such that  $\beta(s)$  represents the probability that  $s$  is the actual world state.

Recall that the environment-agent interactions give rise to a sequence or trajectory:  $s_0, a_0, r_1, s_1, a_1, r_2, s_2, a_2, r_3, \dots$

From the point of view of the partially informed agent, we have a **history**:  $o_0, a_0, r_1, o_1, a_1, r_2, o_2, a_2, r_3, \dots$

The agent needs to maintain its belief based on the current observation.

# Beliefs in the Lab



$$\beta^{o,s}(p) = \frac{\hat{P}(o|p,s) \beta(p)}{\int_{p' \in \mathcal{P}_o} \hat{P}(o|p',s) \beta(p') dp'} \quad (1)$$

where  $\beta(p)$  is the estimated probability that  $p$  is the pose prior to considering observation  $o$ .

# Partially Observable Markov Decision Process (POMDP)



A Partially Observable Markov Decision Process(POMDP) is a tuple  $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma, \Omega, \mathcal{O}, \beta_0 \rangle$  where

- $\mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}$  and  $\gamma$  are as for an MDP.
- $\Omega$  is a set of observations (observation tokens),
- $\mathcal{O}$  is a sensor function specifying the conditional observation probabilities  $\mathcal{O}_{s,a}^o = \mathcal{P}[O_{t+1} = o | S_t = s, A_t = a]$  of receiving observation token  $o \in \Omega$  in state  $s$  after applying action  $a$ <sup>2</sup>.
- $\beta_0$  the initial belief: a probability distribution over the states such that  $\beta_0(s)$  stands for the probability of  $s$  being the true initial state.

---

<sup>2</sup>alternatively:  $\mathcal{O}_s^o = \mathcal{P}[o_t = o | S_t = s]$


# How to model our domain as a POMDP ?



$$\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma, \Omega, \mathcal{O}, b_0 \rangle$$

Is this model appropriate? Limitations? Strengths?





$$\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma, \Omega, \mathcal{O}, b_0 \rangle$$

- State space: Position and orientation of all objects and robot configuration  $\mathcal{S} = \{(x_i, y_i, z_i, rx_i, ry_i, rz_i)\}_{i=1}^{N_{items}} \cup \mathcal{C}_{robot}$
- Action space: pick up or sense if there's object for each point in the workspace  $\mathcal{A} = \{sense(x, y) | x, y \in workspace\} \cup \{pick(x, y) | x, y \in workspace\}$
- Observation Space: whether an item is picked at the last action or whether an object was sensed at the last action  $\{picked, not_picked, sensed, not_sensed\}$



# Optimization Considerations

- Completeness
- Soundness
- Optimality
- Anytime guarantees
- Computational Efficiency
- Generality
- Memory consumption

# Motion Planning

---

# Motion Planning



# Motion Planning

From Oren Salzman's course slides\*

**(Formal) definition of the basic motion-planning\*\* problem**


Let  $\mathcal{R}$  be a robot system with  $d$  degrees of freedom, moving in a known environment cluttered with obstacles. Given start and target configurations  $s$  and  $s'$  for  $\mathcal{R}$ , decide whether there is a collision-free, continuous path  $\tau : [0, 1] \rightarrow \mathcal{C}_{free}$  such that  $\tau(0) = s$  and  $\tau(1) = s'$  and if so, plan such a motion.

---

\*Consider taking course 236901 to learn about algorithms for robotic motion planning.

<https://arxiv.org/pdf/2209.14471.pdf>

# Motion Planning is Hard



More about this soon...

# Control

---

# Example of a Control Problem: Follow the Motion Plan


Given a motion plan  $\tau$ , the trajectory tracking problem is defined by  $(X, U, f, \tau)$  where:

- $X$  is the state space (can be the configuration space and the derivatives in that space (velocities))
- $U$  is the control/input space
- $f: X \times U \rightarrow X$  is the system dynamics what is the next state given the current and the control input, what assumption do we make here?
- $\tau: [0, 1] \rightarrow C_{free}$  is the reference motion plan

The control objective is to find  $u(t)$  such that:

- $\dot{x}(t) = f(x(t), u(t))$  This is the system dynamic equation. Note that it's a ODE
- $\|x(t) - \tau(s(t))\| < \varepsilon$  for some tracking error  $\varepsilon$  We are close enough to the trajectory
- $s(t)$  maps time to path parameter  $[0, 1]$





More about this very soon...

# Summary

## Summary:

- We took a closer look at the structure of a robot
- We examined different models used in AI and robotics

## What next ?

- Cover basic reactive approaches to robotic control

